

Report Common Assignment 3



Counting sort

Rosa Gerardo

Scovotto Luigi

Tortora Francesco

January, 2022

Index

| | |
|---|-----------|
| Index | i |
| 1 Problem description | 1 |
| 1.1 Experimental setup | 1 |
| 1.1.1 Hardware | 2 |
| 1.1.2 Software | 9 |
| 2 Performance | 10 |
| 2.1 Case study n°1 - Global Memory | 11 |
| 2.2 Case study n°2 - Shared Memory | 12 |
| 2.3 Case study n°3 - Texture Memory | 13 |
| 3 Considerations | 14 |
| 3.1 Comparison | 14 |
| 3.2 Premises | 15 |
| 4 How to run | 16 |

Chapter 1

Problem description

Parallelize and Evaluate Performances of "Counting Sort" Algorithm , by using CUDA.

Counting sort is an algorithm for sorting integer numbers. The computational complexity is equal to $O(n)$.

In this report we analyze counting sort algorithm: get execution times, bandwidth and other performances measures when of using:

- Global Memory
- Shared Memory
- Texture Memory

Evaluate performances with different block-grid configurations.

1.1 Experimental setup

All measurements were made using the Google Colab service, with the following setup.

1.1.1 Hardware

CPU

```
processor      : 0
vendor_id     : GenuineIntel
cpu family    : 6
model         : 63
model name    : Intel(R) Xeon(R) CPU @ 2.30GHz
stepping      : 0
microcode     : 0x1
cpu MHz       : 2299.998
cache size    : 46080 KB
physical id   : 0
siblings      : 2
core id       : 0
cpu cores     : 1
apicid        : 0
initial apicid : 0
fpu           : yes
fpu_exception : yes
cpuid level   : 13
wp            : yes
flags         : fpu vme de pse tsc msr pae mce cx8
                apic sep mtrr pge mca cmov pat pse36
                clflush mmx fxsr sse sse2 ss ht syscall
                nx pdpe1gb rdtscp lm constant_tsc rep_good
                nopl xtopology nonstop_tsc cpuid
                tsc_known_freq pni pclmulqdq ssse3
                fma cx16 pcid sse4_1 sse4_2 x2apic
                movbe popcnt aes xsave avx f16c
                rdrand hypervisor lahf_lm abm invpcid_single
```

```

        ssbd ibrs ibpb stibp fsgsbase tsc_adjust
        bml avx2 smep bmi2 erms invpcid xsaveopt
        arat md_clear arch_capabilities
bugs      : cpu_meltdown spectre_v1 spectre_v2
          spec_store_bypass lltf mds swapgs
bogomips  : 4599.99
clflush   size      : 64
cache_alignment : 64
address sizes : 46 bits physical, 48 bits virtual
power management:

processor   : 1
vendor_id   : GenuineIntel
cpu family  : 6
model       : 63
model name  : Intel(R) Xeon(R) CPU @ 2.30GHz
stepping    : 0
microcode   : 0x1
cpu MHz     : 2299.998
cache size  : 46080 KB
physical id : 0
siblings    : 2
core id     : 0
cpu cores   : 1
apicid      : 1
initial apicid : 1
fpu         : yes
fpu_exception : yes
cpuid level : 13
wp          : yes
flags       : fpu vme de pse tsc msr pae mce cx8 apic

```

```

sep mtrr pge mca cmov pat pse36 clflush
mmx fxsr sse sse2 ss ht syscall nx pdpe1gb
rdtscp lm constant_tsc rep_good nopl
xtopology nonstop_tsc cpuid tsc_known_freq
pni pclmulqdq ssse3 fma cx16 pcid sse4_1
sse4_2 x2apic movbe popcnt aes xsave avx
f16c rdrand hypervisor lahf_lm abm
invpcid_single ssbd ibrs ibpb stibp fsgsbase
tsc_adjust bmi1 avx2 smep bmi2 erms invpcid
xsaveopt arat md_clear arch_capabilities
bugs                : cpu_meltdown spectre_v1 spectre_v2
                    spec_store_bypass lltf mds swapgs
bogomips            : 4599.99
clflush size        : 64
cache_alignment     : 64
address sizes       : 46 bits physical, 48 bits virtual
power management:

```

RAM

| | |
|-----------------|-------------|
| MemTotal: | 13302920 kB |
| MemFree: | 236348 kB |
| MemAvailable: | 12350076 kB |
| Buffers: | 283016 kB |
| Cached: | 11566856 kB |
| SwapCached: | 0 kB |
| Active: | 3457004 kB |
| Inactive: | 8911716 kB |
| Active(anon): | 476972 kB |
| Inactive(anon): | 480 kB |
| Active(file): | 2980032 kB |
| Inactive(file): | 8911236 kB |
| Unevictable: | 0 kB |
| Mlocked: | 0 kB |
| SwapTotal: | 0 kB |
| SwapFree: | 0 kB |
| Dirty: | 200 kB |
| Writeback: | 0 kB |
| AnonPages: | 518888 kB |
| Mapped: | 269548 kB |
| Shmem: | 1228 kB |
| KReclaimable: | 548408 kB |
| Slab: | 600016 kB |
| SReclaimable: | 548408 kB |
| SUnreclaim: | 51608 kB |
| KernelStack: | 5904 kB |
| PageTables: | 7172 kB |
| NFS_Unstable: | 0 kB |
| Bounce: | 0 kB |
| WritebackTmp: | 0 kB |

| | |
|------------------|----------------|
| CommitLimit: | 6651460 kB |
| Committed_AS: | 3798912 kB |
| VmallocTotal: | 34359738367 kB |
| VmallocUsed: | 46004 kB |
| VmallocChunk: | 0 kB |
| Percpu: | 1448 kB |
| AnonHugePages: | 0 kB |
| ShmemHugePages: | 0 kB |
| ShmemPmdMapped: | 0 kB |
| FileHugePages: | 0 kB |
| FilePmdMapped: | 0 kB |
| CmaTotal: | 0 kB |
| CmaFree: | 0 kB |
| HugePages_Total: | 0 |
| HugePages_Free: | 0 |
| HugePages_Rsvd: | 0 |
| HugePages_Surp: | 0 |
| Hugepagesize: | 2048 kB |
| Hugetlb: | 0 kB |
| DirectMap4k: | 279360 kB |
| DirectMap2M: | 9154560 kB |
| DirectMap1G: | 6291456 kB |

GPU

Device number: 0

Device name: Tesla K80

Compute capability: 3.7

Clock Rate: 823500 kHz

Total SMs: 13

Shared Memory Per SM: 114688 bytes

Registers Per SM: 131072 32-bit

Max threads per SM: 2048

L2 Cache Size: 1572864 bytes

Total Global Memory: 11996954624 bytes

Memory Clock Rate: 2505000 kHz

Max threads per block: 1024

Max threads in X-dimension of block: 1024

Max threads in Y-dimension of block: 1024

Max threads in Z-dimension of block: 64

Max blocks in X-dimension of grid: 2147483647

Max blocks in Y-dimension of grid: 65535

Max blocks in Z-dimension of grid: 65535

Shared Memory Per Block: 49152 bytes

Registers Per Block: 65536 32-bit

Warp size: 32

BANDWIDTH

Device 0: Tesla K80

Range Mode

Host to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

| Transfer Size (Bytes) | Bandwidth(MB/s) |
|-----------------------|-----------------|
| 1000 | 278.7 |
| 101000 | 6240.3 |
| 201000 | 6613.4 |
| 301000 | 6923.8 |
| 401000 | 7064.9 |
| 501000 | 7197.9 |
| 601000 | 7046.2 |
| 701000 | 7217.7 |
| 801000 | 7333.4 |
| 901000 | 7454.2 |

Device to Host Bandwidth, 1 Device(s)

PINNED Memory Transfers

| Transfer Size (Bytes) | Bandwidth(MB/s) |
|-----------------------|-----------------|
| 1000 | 419.3 |
| 101000 | 6563.5 |
| 201000 | 7137.2 |
| 301000 | 7351.7 |
| 401000 | 7361.0 |
| 501000 | 7443.2 |
| 601000 | 7505.5 |
| 701000 | 7519.7 |
| 801000 | 7579.2 |
| 901000 | 7571.1 |

Device to Device Bandwidth, 1 Device(s)

PINNED Memory Transfers

| Transfer Size (Bytes) | Bandwidth(MB/s) |
|-----------------------|-----------------|
| 1000 | 251.6 |
| 101000 | 25527.1 |
| 201000 | 44132.9 |
| 301000 | 57341.3 |
| 401000 | 75400.8 |
| 501000 | 83089.5 |
| 601000 | 89331.3 |
| 701000 | 102785.9 |
| 801000 | 91342.3 |
| 901000 | 80354.0 |

1.1.2 Software

- NVIDIA-SMI 510.39.01
- Driver Version: 460.32.03
- CUDA Version: 11.2
- Ubuntu 18.04.5 LTS

Chapter 2


Performance

2.1 Case study n°1 - Global Memory

In this case study, we have analyzed the counting sort algorithm using global memory. The size of the problem is 102 400 000 and than we evaluate the performances with block size of:

- 128: with this choice we have that $2048/128 = 16$ blocks. The occupancy is 100%.
- 256: with this choice we have that $2048/256 = 8$ blocks. The occupancy is 100%.
- 512: with this choice we have that $2048/512 = 4$ blocks. The occupancy is 100%.
- 1024: with this choice we have that $2048/1024 = 2$ blocks. The occupancy is 100%.

| N | BlockSize | GridSize | gips | time_sec |
|-------------|-----------|----------|------------|------------|
| 102 400 000 | 128 | 800 000 | 40 278 885 | 0,04830324 |
| 102 400 000 | 256 | 400 000 | 40 307 517 | 0,04826894 |
| 102 400 000 | 512 | 200 000 | 40 325 467 | 0,04824748 |
| 102 400 000 | 1024 | 100 000 | 40 344 692 | 0,04822454 |

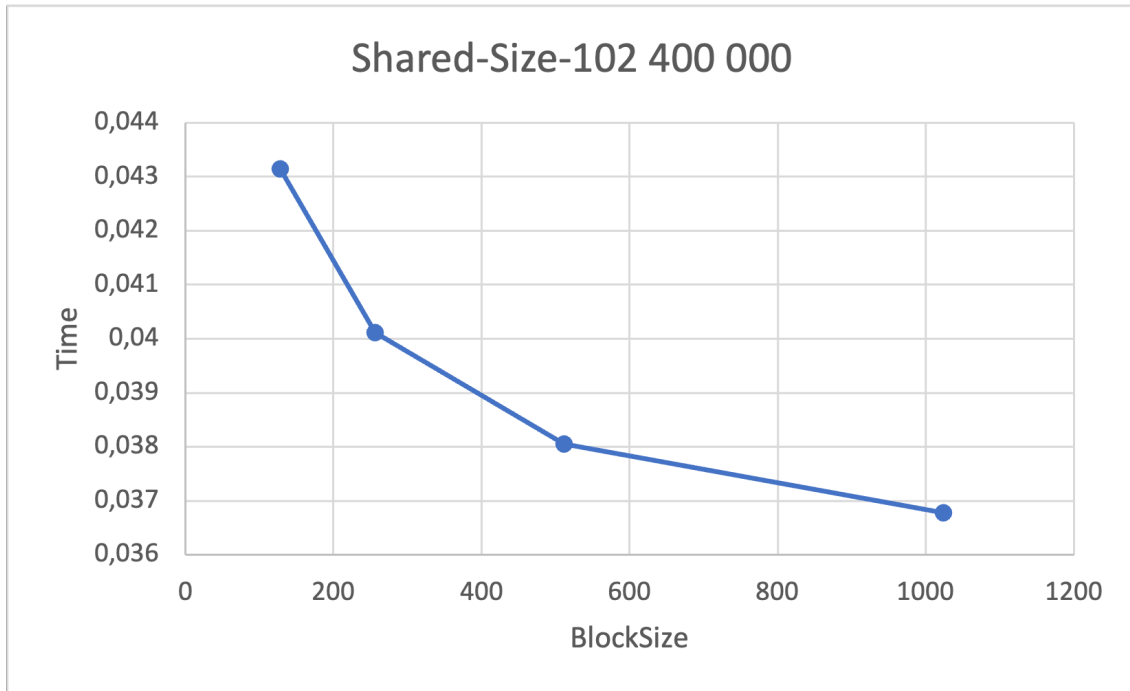


2.2 Case study n°2 - Shared Memory

In this case study, we have analyzed the counting sort algorithm using shared memory. The size of the problem is 102 400 000 and than we evaluate the performances with block size of:

- 128: with this choice we have that $2048/128 = 16$ blocks. The occupancy is 100%.
- 256: with this choice we have that $2048/256 = 8$ blocks. The occupancy is 100%.
- 512: with this choice we have that $2048/512 = 4$ blocks. The occupancy is 100%.
- 1024: with this choice we have that $2048/1024 = 2$ blocks. The occupancy is 100%.

| N | BlockSize | GridSize | gips | time_sec |
|-------------|-----------|----------|------------|------------|
| 102 400 000 | 128 | 800 000 | 52 675 126 | 0,04314758 |
| 102 400 000 | 256 | 400 000 | 57 151 101 | 0,04011825 |
| 102 400 000 | 512 | 200 000 | 61 304 339 | 0,03805279 |
| 102 400 000 | 1024 | 100 000 | 65 612 982 | 0,03677322 |

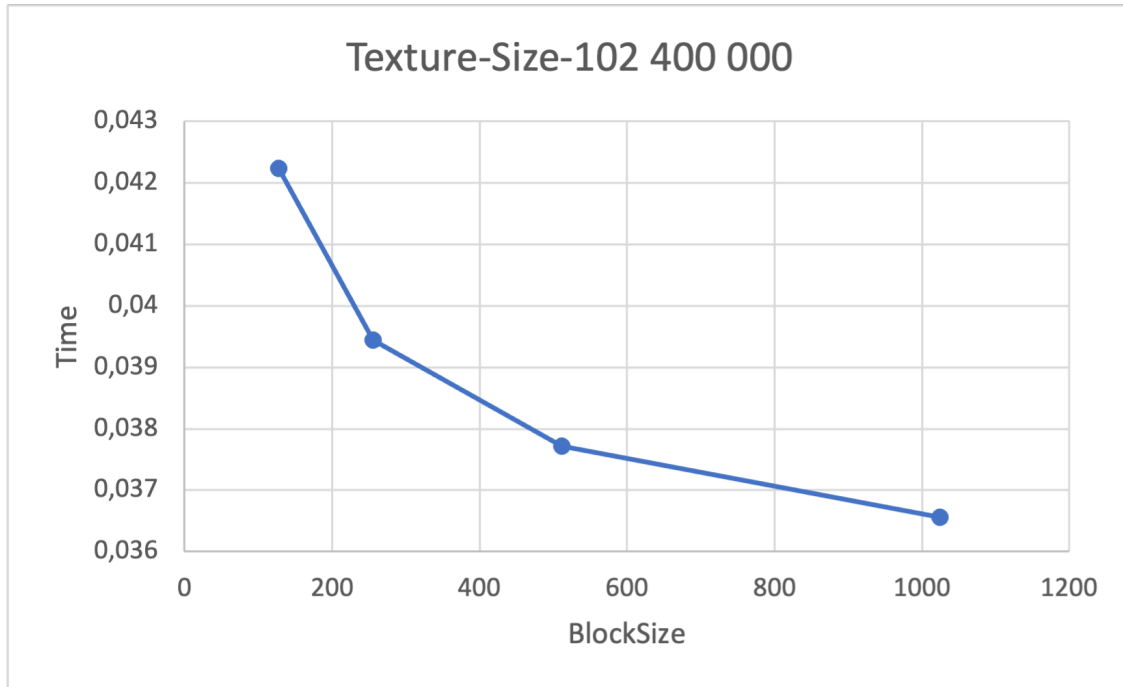


2.3 Case study n°3 - Texture Memory

In this case study, we have analyzed the counting sort algorithm using texture memory. The size of the problem is 102 400 000 and than we evaluate the performances with block size of:

- 128: with this choice we have that $2048/128 = 16$ blocks. The occupancy is 100%.
- 256: with this choice we have that $2048/256 = 8$ blocks. The occupancy is 100%.
- 512: with this choice we have that $2048/512 = 4$ blocks. The occupancy is 100%.
- 1024: with this choice we have that $2048/1024 = 2$ blocks. The occupancy is 100%.


| N | BlockSize | GridSize | gips | time_sec |
|-------------|-----------|----------|------------|------------|
| 102 400 000 | 128 | 800 000 | 53 808 727 | 0,04223858 |
| 102 400 000 | 256 | 400 000 | 58 131 895 | 0,03944138 |
| 102 400 000 | 512 | 200 000 | 61 855 504 | 0,03771372 |
| 102 400 000 | 1024 | 100 000 | 65 988 295 | 0,03656407 |



Chapter 3

Considerations

3.1 Comparison



As can be seen from the graph above, shared and textured versions are influenced by the BlockSize: as the size increases, the time decreases.

Comparing the various versions, a clear improvement can be seen in the use of shared com-

pared to global as the shared has a very low latency. In the texture, there is an improvement over the global and shared versions, but only a small improvement in terms of time compared to the shared.

3.2 Premises

- Global memory is L1/L2-cache and texture memory is a separate texture cache.
- L2 is global so each access to global memory goes through L2 first, L1 is local to the multiprocessor.
- Texture memory is read-only, our program cannot dictate what is in the texture cache, so it is transparent to the program. (Behaves much like a CPU cache).
- We used atomic functions because the counting sort algorithm required that several threads access the same locations but above all it avoids data inconsistency due to the race condition.

Referring to the documentation, we realized that the atomic function was useful for our purpose because:

an atomic function performs a read-modify-write atomic operation on one 32-bit or 64-bit word residing in global or shared memory.

The operation is atomic in the sense that it is guaranteed to be performed without interference from other threads. In other words, no other thread can access this address until the operation is complete.

Atomic functions do not act as memory fences and do not imply synchronization or ordering constraints for memory operations.

In particular, we used `atomicAdd` and `atomicSub` function, the following URLs are referred to the specified functions documentation:

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomicadd>

<https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#atomicsub>

Chapter 4

How to run

Use the colab notebook in the folder or go to this URL:

<https://drive.google.com/drive/u/1/folders/10ZMNoKnH-foAqqGiVjpBPBVHRx-Z5TAw>

You can find the complete project on GitHub: <https://github.com/scov8/CommonAssignment3>

The starting point for the global and shared counting sort was taken here:

<https://github.com/avtrunov/CUDA/blob/500bf0f1ca9af066f8dba6572d7b11121046d3a6/CountingSort.cu>

https://github.com/jtruong27/ECE413/blob/9713c61a9c5dc10c77b5d79def7ea83abc21fcb0/Homework_5/CountingSort_CUDA/counting_sort_kernel.cu

The previous year's group 02 files proposed by the professor during the course were used for file generation and extraction.

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. To view a copy of this license, visit

<http://creativecommons.org/licenses/by-nc-sa/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA.